

# X: A PROTOCOL FOR DECENTRALIZED EXCHANGE WITH LIMIT ORDERS

G.P. SAGGESE AND PAUL SMITH

## CONTENTS

1. Introduction	2
1.1. Limit orders as an interface	2
1.2. How exchanges create value	2
1.3. The advantages of X	3
2. Matching liquidity	3
2.1. Breaking the symmetry	4
2.2. Basket case	4
3. Protocol introduction	4
3.1. Base/quote tokens	4
3.2. Order attributes and notation	5
3.3. Limit order as a set of inequalities	6
3.4. Order normalization	7
4. Reference clearing prices and prioritization	8
4.1. Lit exchange bid-ask midpoint reference price	8
4.2. Order crossing	8
5. Formulation of the general DaoSwap problem	9
5.1. Limit order inequalities	9
5.2. No arbitrage	9
5.3. Conservation of exchanged value	10
5.4. Conservation of exchanged tokens	10
5.5. Combining the constraints	10
5.6. Solution of the general DaoSwap problem	10
5.7. Reformulating the DaoCross problem	10
6. Alternative solutions for token swapping	10
6.1. DaoSwap	11
6.2. DaoCross	11
6.3. Limit Order Book (LOB)	12
6.4. Automatic Market Makers (AMM)	12
References	13

---

*Date:* April 6, 2023.

With contributions from Samarth KaPatel, Grisha Pomazkin, Juraj Smeriga, Danya Tikhomirov, Nina Trubacheva, Dan Yachmenev, Vladimir Yakovenko, and Tamara Zhordaniya.

## 1. INTRODUCTION

X is a new protocol that enables decentralized token exchange using limit orders. It brings together the advantages of efficiency and ease-of-use of centralized exchanges with the transparency and self-custody of decentralized approaches.

**1.1. Limit orders as an interface.** A limit order expresses a commitment to buy or to sell an asset, up to a certain quantity, at a price that is as good or better than the limit price. As supply and demand are represented by quantity available at a given price, limit orders provide a way for buyers and sellers to participate in price discovery and exchange.

We adopt the perspective that limit orders may be thought of as an interface to market participation. That is, limit orders provide participants with a means to engage with markets. We further promote this interface as one that is natural, flexible, and elegant. It is natural in that it expresses supply or demand directly in terms of quantity and price, with a clearly defined maximum possible exposure in a commitment to trade. The interface is flexible because it allows participants to construct personalized supply and demand curves by combining multiple limit orders. The elegance follows from the simplicity and ease-of-use of the interface.

In any case, limit orders and their aggregation into limit order books (LOBs) are tremendously popular. They are ubiquitous in traditional financial markets and dominate cryptocurrency trading on centralized exchanges. Interestingly, however, limit orders currently do not play as prominent a role in the realm of decentralized finance. Instead, alternative approaches that rely upon automated market makers (AMMs) have taken center stage, and these approaches offer a different set of advantages and disadvantages.

**1.2. How exchanges create value.** To better understand the comparative advantages and disadvantages of different approaches to token exchange, we step back and consider the purpose and functions of exchanges. Exchanges create value by bringing together different types of participants, such as investors, traders, hedgers, brokers, arbitrageurs, and market makers. By providing a common meeting ground, exchanges facilitate matching supply and demand.

Typically, an exchange (or swap) requires two components:

- (1) Trade matching: participants find a counterparty agreeing on quantity and price of the assets to swap (this includes price discovery)
- (2) Trade settlement: the assets are actually swapped after finding matching counterparties

In decentralized finance a token swap can be accomplished in different ways, e.g., using

- a centralized exchange (e.g., Coinbase, Binance, OkX)
- a decentralized exchange (e.g., Uniswap, SushiSwap)

Centralized exchanges are usually organized around different versions of a central limit order book (CLOB), whereas decentralized exchanges have been organized as

- off-chain order books (e.g., 1inch)
- on-chain order books (e.g., SwapSwap, KyberSwap)
- automataged market makers (AMM) (e.g., Uniswap)

Any off-chain solution runs into custodial and censorship issues that stand in direct conflict with the ethos of decentralization and self-determination.

Off-chain order books are an hybrid version of CEX and DEX where the matching

On-chain order book matching has the advantages of being custodial, since users are completely in charge of their funds. A major disadvantage of using on-chain limit order books in decentralized exchange is cost of computation. Continual submission and cancelation of orders incurs costs, as does the ongoing process of matching supply and demand.

Exchanges are compensated for this value creation by collecting transaction fees.

**1.3. The advantages of X.** *X* retains the key advantages of on-chain decentralized exchange while overcoming the issues of cost. Additionally, it includes an implementation that preserves the LOB interface but deepens the pool of available liquidity in matching supply and demand.

*X* addresses cost issues primarily in two ways: (1) discretizing time; and (2) scaling with layer 2 solutions. By addressing these issues, we combine the novel advantages of decentralized exchange with the utility and ease-of-use of the familiar interface of limit orders.

Behind the interface, we propose two smart contracts<sup>1</sup> for ERC-20<sup>2</sup> token exchange that match supply and demand differently depending upon whether there exists an external reference price or whether the contract also performs the role of price discovery.

DaoCross is a decentralized liquidity pool that crosses buy and sell orders with respect to an external reference price, i.e., a price oracle. This allows traders to interact directly with each other and share price improvement with respect to exchanges, e.g., by trading at bid-ask midpoint. Additionally, order intentions are not publicly disclosed prior to a cross, which enables block or other high-volume traders to execute quickly with minimal market impact.

DaoSwap performs price discovery by matching supply and demand at regular time intervals. In the simplest case, the clearing price is determined by a Walrasian-style auction ([12]). In cases where liquidity is pooled more broadly, the auction clearing mechanism satisfies a collection of conditional inequalities.

While DaoSwap and DaoCross differ in how the token exchange rate is determined, both share several common advantages:

- (1) Low cost: they are a DeFi primitive that allows trading tokens peer-to-peer without incurring spread costs
- (2) Capital efficiency: participants provide liquidity only when they wish to trade, and only in the amount they wish to trade
- (3) No impermanent loss: liquidity providers are not exposed to the risk of impermanent loss (a form of unrealized loss that becomes realized upon withdrawing liquidity), unlike the case with other decentralized exchange protocols based on automated market makers
- (4) No intermediary custody: exchanged tokens always remain under the control of their respective owners
- (5) No rent extraction from high-frequency traders: speed alone does not confer an advantage to traders, as the discrete timing prevents predatory tactics such as front-running, limit order scalping, and spoofing, which are known issues seen with continuous limit order books
- (6) Ease-of-use: the interaction between buyers and sellers occurs through smart contracts, with a website front-end available

## 2. MATCHING LIQUIDITY

Suppose there are two parties, Alice an **ETH** holder and Bob a **wBTC** holder (wrapped bitcoin, an ERC-20 compliant coin that tracks bitcoin), who wish to swap tokens at a competitive rate. Suppose also that there are many additional **ETH** and **wBTC** holders who may be interested in participating in a swap. How should the liquidity be pooled? At first glance, it appears that there should be a single pool of liquidity. The wrinkle manifests in determining how to express a desire to trade. A simple expression of a desire to trade is a limit order representing a commitment to trade up to  $q$  in quantity of a token at a token exchange rate up to  $p$ . But what if some parties express quantity in terms **ETH** and some in terms of **wBTC**, and some parties express limit prices in terms of **ETH** per **wBTC**, while others express limit prices in terms of **wBTC** per **ETH**? Under these conditions, setting

---

<sup>1</sup><https://ethereum.org/en/developers/docs/smart-contracts/>

<sup>2</sup><https://ethereum.org/en/developers/docs/standards/tokens/erc-20/>

a single clearing price and determining fill prioritization rules raises several questions. We will consider and address these in the sequel, but to begin, we will discuss a more constrained setting.

**2.1. Breaking the symmetry.** To simplify, we follow the practice of FX (foreign exchange) futures markets, which break the symmetry by specifying two non-interchangeable roles:

- (1) a *base* currency (for quantity)
- (2) a *quote* currency (for price)

The roles are expressed by writing *base currency/quote currency*, e.g., “EUR/USD”.

A limit order then consists of a quantity expressed in the units of the base currency and a price expressed in terms of the quote currency (per unit or suitable multiple of base currency). See, for example, [6] (e.g., footnotes on page 25) for usage of this terminology. See [7] for how this works in practice for Euro FX contracts, where contract units are denominated in Euros and price is quoted in U.S. dollars and cents per Euro increment.

By breaking the symmetry, one may run a standard limit order book, hold a classical Walrasian-style auction, and handle size quantization effects by using one of a variety of priority rules based on quantity, price, and time. See [4] for discussion around priority rule variations and their effects on market quality outcomes.

An effect of breaking the symmetry in this way is that either liquidity for a pair of currencies may be expressed in only one form, or liquidity for a pair of currencies is split across two separate contracts (with roles of base and quote token reversed), each with its own limit order book. Using our example, this would mean treating  $\text{wBTC/ETH}$  and  $\text{ETH/wBTC}$  as separate token pairs, even though the union of the underlying parties and sources of liquidity are the same.

**2.2. Basket case.** It seems plausible that by retaining the symmetry in the two-token case and instead treating a pair of tokens as a single source of liquidity, one may contrive a more efficient exchange mechanism.

This minor expansion, however, suggests widening even further the universe of eligible swaps. For example, if there are three tokens available for exchange, one could contemplate many-for-one, one-for-many, or many-for-many token swaps, all to be carried out in an atomic fashion, and perhaps with complex constraints. A toy case along these lines one may consider is as follows:

**Problem 2.2.1** (Triangular liquidity). Suppose there are the following three parties:

- Party A, who holds  $\text{wBTC}$  and wants  $\text{ETH}$
- Party B, who holds  $\text{ETH}$  and wants  $\text{DAI}$
- Party C, who holds  $\text{DAI}$  and wants  $\text{wBTC}$

How should an auction be structured to “best” facilitate exchange?

While there may be other use cases and even demand for such auctions, unfortunately, the problem in general is NP-complete (e.g., [13]), which is a significant limiting factor in terms of feasibility and auction expense.

### 3. PROTOCOL INTRODUCTION

In the previous section, we discussed the notions of base currency and quote currency from foreign exchange. Here we extend them to tokens.

**3.1. Base/quote tokens.** A token (ordered) pair consists of a *base token* and a *quote token*. The shorthand notation for expressing the pair is *base token/quote token* (e.g., “ $\text{wBTC/ETH}$ ”), and as such indicates the respective roles of the tokens.

By convention, *quantity* to exchange is expressed in terms of the base token, and *price* is a token exchange rate expressed in terms of the quote token per unit of base token (or multiple thereof).

*Example 3.1.1* (base/quote tokens). By way of example, in the pair `wBTC/ETH`, `wBTC` is the base token and `ETH` is the quote token. Interest to trade is expressed in terms of orders with `wBTC` as base token and `ETH` as quote token, as follows:

- A “buy” order represents a commitment to purchase the base token `wBTC` and pay with the quote token `ETH`
- A “sell” order represents a commitment to sell the base token `wBTC` and receive the quote token `ETH`.
- Quantity  $q$  is in terms of the base token (“buy/sell a certain amount of `wBTC`”)
- Limit price is in terms of the quote token (“buy/sell `wBTC` with a limit price of  $p$  `ETH` per `wBTC`”)

A party who places a “buy” order must have the quote token (`ETH`) in custody, i.e., in its wallet. A party who places a “sell” order must have the base token (`wBTC`) in custody.

**3.2. Order attributes and notation.** We introduce the following tuple notation for a general limit order (valid for both `DaoCross` and `DaoSwap`).

**Definition 3.2.1** (Limit order). A `DaoCross` or `DaoSwap` limit order is represented by a tuple of the form

```
(timestamp,
  action,
  quantity,
  base_token,
  limit_price,
  quote_token,
  deposit_address)
```

The quantities are arranged to make the order simple to read in natural language: “At timestamp `timestamp` create an order to `action` a `quantity` of the token `base_token` for a limit price of `limit_price` with respect to the token `quote_token` and deposit the resulting tokens at `deposit_address`.”

We have previously discussed the roles of `action`, `quantity`, `base_token`, `limit_price`, and `quote_token`. The roles of `timestamp` include determining eligibility in a swap or cross and can extend to influencing order fill priority. The `deposit_address` attribute mirrors standard functionality available in traditional finance, e.g., in purchasing T-Bills on TreasuryDirect<sup>3</sup>. This feature facilitates account management through the use of multiple wallets. For example, a miner may have a supply of `wBTC` collected and held in one wallet which it would like to systematically diversify into `ETH` and perhaps other tokens. In specifying a deposit address, it may use a separate wallet to collect incoming `ETH`.

For brevity, in the sequel we may:

- omit `timestamp` and `deposit_address` when not relevant to the discussion
- omit the (infinite) `limit_price` for market orders

*Example 3.2.1* (Limit order notation). The order

```
(1678660406, buy, 3.2, ETH, 4.0, wBTC, 0xdeadcd0e)
```

corresponds to the natural language description: “At timestamp `Mon Mar 13 2023 02:33:25 GMT+0000`, the user commits to buy up to 3.2 units of `ETH` in exchange for `wBTC` up to a `limit_price` of 4.0 `wBTC` per `ETH` with proceeds deposited at `0xdeadcd0e`”.

<sup>3</sup><https://www.treasurydirect.gov>

Next we introduce some examples of market order and limit order behavior when a clearing rate has been set. In particular, consider a swap for the tokens ETH, wBTC and assume that the exchange rate between ETH and wBTC is fixed at 0.2 (i.e., 0.2 wBTC can be exchanged for 1 ETH and vice versa).

*Example 3.2.2* (Market order notation and clearing). The following market orders omit `timestamp`, `limit_price`, and `deposit_address` in favor of emphasizing the amounts of tokens exchanged:

- An order (buy, 1.0, ETH, wBTC) means buying 1 ETH in exchange for 0.2 wBTC
- An order (sell, 1.0, ETH, wBTC) means selling 1 ETH, receiving the corresponding amount of 0.2 wBTC
- An order (buy, 1.0, wBTC, ETH) means buying 1 wBTC, paying with 5 ETH
- An order (sell, 1.0, wBTC, ETH) means exchanging 1 wBTC in return for 5 ETH

Next we consider the behavior of limit orders under the same prevailing exchange rate (and we assume that there is sufficient supply of tokens to fully fill the orders).

*Example 3.2.3* (Executable buy limit order). A limit order

(buy, 1.0, ETH, 0.5, wBTC)

means “buy up to 1 ETH in exchange for wBTC at a rate up to 0.5 wBTC per ETH.” In this case, since the price of one ETH is equal to 0.2 wBTC, the order can be executed at the prevailing market rate.

*Example 3.2.4* (Non-executable buy limit order). On the other hand, a limit order

(buy, 1.0, ETH, 0.1, wBTC)

requires that the rate of wBTC per ETH be lower than the current market rate, and so the limit price prevents a token swap from being carried out.

*Example 3.2.5* (Non-executable sell limit order). A limit order

(sell, 1.0, ETH, 0.5, wBTC)

means “sell up to 1 unit of ETH in exchange for wBTC at a rate down to 0.5 wBTC per ETH.” Since the current rate of wBTC per ETH is 0.2, which is below the limit price of 0.5, the order cannot be executed.

**3.3. Limit order as a set of inequalities.** Any limit order as defined above can be translated into inequalities involving quantity of exchanged tokens and token exchange rates. Multiple limit orders can be converted into a system of inequalities, which collectively constrain potential exchange outcomes.

In both DaoCross and DaoSwap, a swap between a given base/quote token pair occurs at a single exchange rate, i.e., the clearing price is the same for all executed orders. On the other hand, quantity exchanged is specific to each order and each order’s constraint on quantity exchanged cannot be violated. In the same way each order can be executed with a non-null quantity only when the single exchange rate is compatible with the limit price.

In the sequel we use an asterisk to denote clearing quantity and exchange rate:

- $q_{base}^*(o_i)$  denotes quantity of the base token exchanged in a swap from order  $o_i$
- $p_{quote\_per\_base}^*$  denotes exchange rate between the quote and base tokens for all the swaps

Note that

$$p_{quote\_per\_base} = \frac{1}{p_{base\_per\_quote}}.$$

If prices of base and quote tokens are expressed in terms of a common currency (e.g., USD) then it holds that:

$$p_{quote\_per\_base} = \frac{p_{quote}}{p_{base}},$$

where the numeraire plays the role of a quote token for both  $p_{quote}$  and  $p_{base}$  and is implicit.

Given an order  $o_i = (\text{buy}, q, A, p, B)$  we indicate with:

- $action(o_i)$  the desired action (**buy** or **sell**)
- $q_{base}(o_i)$  the maximum desired quantity  $q$
- $base(o_i)$  the base token  $A$
- $p(o_i)$  the limit price  $p$  (in terms of quote per base)
- $quote(o_i)$  the quote token  $B$

**3.3.1. Inequalities for buy order.** An order of the form  $o_i = (\text{buy}, q, A, p, B)$  means “buy  $q$  units of token  $A$  in exchange for token  $B$  with a limit price up to  $p$   $B$  per unit of  $A$ ”, and it corresponds to the following constraint on realized quantity exchanged and clearing exchange rate:

$$(p_{B\_per\_A}^* \leq p(o_i)) \wedge (0 \leq q_A^*(o_i) \leq q(o_i)) \vee (q_A^*(o_i) = 0)$$

The constraint on the quantity exchanged is conditioned on the corresponding price satisfying the desired limit price constraint: if the desired limit price constraint is not met, the swap cannot be carried out and the exchanged quantity is 0.

*Example 3.3.1* (Inequalities for buy order). An order of the form  $o_1 = (\text{buy}, 2, A, 3, B)$  means “buy 2 units of token  $A$  in exchange for token  $B$  with a limit price up to 3  $B$  per unit of  $A$ ”, and it corresponds to the following constraint on realized quantity exchanged and clearing exchange rate:

$$(p_{B\_per\_A}^* \leq 3) \wedge (0 \leq q_A^*(o_1) \leq 2) \vee (q_A^*(o_1) = 0)$$

*Example 3.3.2* (Inequality for buy market order). A market order (i.e., without a limit price) has no constraint on exchange rate and thus is reduced to

$$0 \leq q_A^*(o_1) \leq q(o_i)$$

**3.3.2. Inequalities for sell order.** An order of the form  $o_i = (\text{sell}, q, A, p, B)$  means “sell  $q$  units of token  $A$  in exchange for token  $B$  with a limit price of at least  $p$   $B$  per unit of  $A$ ”, and it corresponds to the following constraint on realized quantity exchanged and clearing exchange rate:

$$(p_{B\_per\_A}^* \geq p(o_i)) \wedge (0 \leq q_A^*(o_i) \leq q(o_i)) \vee (q_A^*(o_i) = 0)$$

*Example 3.3.3* (Inequalities for sell order). In the same way, an order like  $o_2 = (\text{sell}, 3, A, 2, B)$  means “sell 3 units of token  $A$  in exchange for token  $B$  with a limit price of at least 2  $B$  per unit of  $A$ ”, which corresponds to

$$(p_{B\_per\_A}^*(o_2) \geq 2) \wedge (0 \leq q_A^*(o_2) \leq 3) \vee (q_A^*(o_2) = 0)$$

**3.4. Order normalization.** When the exchange rate between two tokens is known, it is possible to convert buy/sell orders for both  $A$  and  $B$  tokens into buy/sell orders where all orders have token  $A$  only, according to the transformation below.

Let

$$o_1 = (\text{buy}, q, A, p_{B\_per\_A}, B).$$

Suppose that  $p_{B\_per\_A}^*$  is fixed and known. Then the order

$$o_2 = (\text{sell}, q', B, 1/p_{B\_per\_A}, A)$$

may be handled in order crossing in the same way that  $o_1$  may be provided that

$$q' = q \cdot p_{B\_per\_A}^*$$

#### 4. REFERENCE CLEARING PRICES AND PRIORITIZATION

DaoCross relies on a *price oracle* for determining the effective token exchange rate in a cross. The price oracle may come from a lit exchange, centralized or decentralized, or even on-chain automated market makers such as Uniswap ([2, §2.2]).

The case of using an automated market maker as a price oracle is relatively straightforward, provided there is an automated market maker trading the target currency pair with sufficient liquidity.

To use an exchange as a price reference, we must consider some additional steps. First, exchanges typically use a dollar stablecoin (e.g., USDC, USDT) as the quote currency and all other currencies as base currencies. Our example of BTC/ETH would be considered a cross pair in the world of traditional finance. Because most cross pairs are not traded directly on exchanges, we must derive a clearing price from pairs involving stablecoins.

**4.1. Lit exchange bid-ask midpoint reference price.** Let  $\text{bid}_{\text{BTC}}$ ,  $\text{ask}_{\text{BTC}}$  be stablecoin-denominated top-of-book bid-ask prices for  $\text{wBTC}$  (e.g., expressed in USDC), and  $\text{bid}_{\text{ETH}}$ ,  $\text{ask}_{\text{ETH}}$  be the analogous prices for ETH. Then, a commitment to buy one  $\text{wBTC}$  and pay ETH would clear at a midpoint price of

$$\text{wBTC/ETH}_{\text{midpoint}} = \frac{\text{bid}_{\text{BTC}} + \text{ask}_{\text{BTC}}}{\text{bid}_{\text{ETH}} + \text{ask}_{\text{ETH}}}$$

expressed in BTC per ETH. Note that if the dollar price of BTH is significantly more than the dollar price of ETH, then this price implies paying many multiples of ETH for BTC (as is the case at the time of this writing).

In general, DaoCross may use the following reference price for a base/quote token pair, where the bids and asks come from a lit exchange and are expressed in terms of stablecoin prices:

$$\frac{\text{bid}_{\text{quote token}} + \text{ask}_{\text{quote token}}}{\text{bid}_{\text{base token}} + \text{ask}_{\text{base token}}} \quad (1)$$

When using either an on-chain price oracle or an exchange as a price oracle, it is possible to perform a time or volume weighted average of prices so as to avoid extreme price variations and to mitigate the risk and effects of any market manipulation attempts.

**4.2. Order crossing.** At regular time intervals, order submissions are cut off and reference prices are determined. An element of randomness is used to determine order cutoff times and reference price cutoff times in order to mitigate manipulative behaviors.

An order is eligible for matching if its timestamp is within the cutoff window and if the external reference price does not exceed its limit price.

Except on occasions where eligible buy/sell volume is perfectly matched, not all orders can be fully crossed. There are also discretization effects to consider arising from discrete order sizes and a discrete price grid [5][§3.2.1]. See [4] for a discussion of the trade-offs surrounding different prioritization rule choices. See [11] for the prioritization rules used by one of Morgan Stanley’s equity liquidity pools.

**4.2.1. Priority rules.** DaoCross prioritizes fills according to:

- Volume (higher volume comes first in priority)
- Price (higher limit price breaks volume ties)
- Timestamp (earlier timestamp breaks ties in volume and price)

If, in the unlikely scenario that all three of these parameters perfectly agree, a certain priority is not guaranteed.



4.2.2. *Matching algorithm.* The mechanism for matching eligible buy and sell orders consists of two priority queues, one for eligible buy orders and one for eligible sell orders. Priority is determined according to the volume/price/timestamp priority rules introduced above. Top-of-queue orders are compared, and the lesser of the two volumes is fully filled. Once an order is fully filled at the established reference price. Once an order is fully filled, it is removed from the priority queue. The procedure continues until one of the two priority queues is empty.

4.2.3. *Computational complexity.* Let  $n$  denote the sum (or max) of the number of eligible buy and sell orders. Priority queue construction occurs in  $O(n)$  time, order removal costs  $O(n \log n)$ , and order remove occurs  $O(n)$  times. So, the computational time complexity of the task is  $O(n \log n)$ . Memory requirements are  $O(n)$ .

4.2.4. *DaoSwap variant.* The mechanics for DaoSwap are similar to those above, with the primary difference being that an auction is used to determine a single clearing price. Variations in prioritization rules also possible.

## 5. FORMULATION OF THE GENERAL DAO-SWAP PROBLEM

The general problem of determining the token equilibrium price and the allocation among the swap participants can be formulated in terms of the inequalities on quantities and prices corresponding to the orders and on the need that the total quantity exchanged of each token be preserved across the swaps, i.e., the quantity bought for each token is equal to the quantity sold.

5.1. **Limit order inequalities.** Let  $\{o_i\}_{i=1}^n$  be a set of limit orders, each of the form

$$o_i = (a_i, q_i, \text{base\_token}_i, p_i, \text{quote\_token}_i)$$

where the quote token is implicit in the quantity  $q_i$  and the quote and base tokens are implicit in the limit price  $p_i$ . To encode directionality in the limit price inequality, we introduce

$$\mathcal{I}(a) := \begin{cases} +1 & \text{for } a = \text{buy} \\ -1 & \text{for } a = \text{sell} \end{cases}$$

For each  $i$ , define

$$c_i := \left( p_{\text{quote\_token}_i \text{ per } \text{base\_token}_i}^* \leq \mathcal{I}(a_i) \cdot p_i \right)$$

where  $p_{\text{quote\_token}_i \text{ per } \text{base\_token}_i}^*$  is the unique clearing price for the indicated base token/quote token ordered pair. TODO(Paul): revive the “to token” notation. Then, each limit order  $o_i$  may be expressed as

$$c_i \wedge (q_i^* \leq q_i) \vee (q_i^* = 0)$$

5.2. **No arbitrage.** Let  $T$  denote the intersection of the union of all base tokens and the union of all quote tokens:

$$T := (\cup_i \text{base\_token}_i) \cap (\cup_i \text{quote\_token}_i)$$

The set of tokens  $T$  is the set of tokens eligible for swapping. We require the following *no arbitrage* constraints be satisfied for all tokens  $u, v, w \in T$  involved in a swap:

$$p_{u \text{ per } v}^* \cdot p_{v \text{ per } w}^* = p_{u \text{ per } w}^*$$

Note that, in the case  $w = u$ , this enforces a unique exchange rate between a pair of tokens (changing the roles of the base and quote tokens inverts the exchange rate).

**5.3. Conservation of exchanged value.** For each order  $i$ , let  $\tilde{q}_i^*$  denote the quantity of `quote_tokeni` exchanged. Then we must have

$$\tilde{q}_i^* = q_i^* \cdot p_{\text{quote\_token}_i \text{ per base\_token}_i}^*$$

for each  $i$ . In effect, this enforces the market clearing price for each order.

**5.4. Conservation of exchanged tokens.** Assume by convention that a buy order removes base tokens from the system and provides quote tokens to the system, whereas a sell order does the opposite.

For each token  $u \in T$ , define the indicator function  $\mathcal{T} : T \rightarrow \{0, 1\}$  via

$$\mathcal{T}_u(v) = \begin{cases} 1 & \text{if } v = u \\ 0 & \text{if } v \neq u \end{cases}$$

Then, for each  $u \in T$ , the following conservation law must hold:

$$\sum_i \mathcal{T}_u(\text{base\_token}_i) \cdot q_i^* = - \sum_i \mathcal{T}_u(\text{quote\_token}_i) \cdot \tilde{q}_i^*$$

Note that this translates into one equality per token participating in the swap. Note also that this is a global constraint on the swap rather than local to each order; in particular, it expresses the notion that each filled order must have, across the collection of orders, suitable counterparties.

**5.5. Combining the constraints.** We now collect the constraints that must be satisfied:

$$\begin{cases} c_i \wedge (q_i^* \leq q_i) \vee (q_i^* = 0) & \forall i \in \{1, \dots, n\} \\ \tilde{q}_i^* = q_i^* \cdot p_{\text{quote\_token}_i \text{ per base\_token}_i}^* & \forall i \in \{1, \dots, n\} \\ p_{u \text{ per } v}^* \cdot p_{v \text{ per } w}^* = p_{u \text{ per } w}^* & \forall u, v, w \in T \\ \sum_{j=1}^n \mathcal{T}_u(\text{base\_token}_j) \cdot q_j^* = - \sum_{j=1}^n \mathcal{T}_u(\text{quote\_token}_j) \cdot \tilde{q}_j^* & \forall u \in T \end{cases} \quad (2)$$

**5.6. Solution of the general DaoSwap problem.** The general solution of the DaoSwap problem must satisfy a set of non-linear, combinatorial constraints, whose solution is NP-complete.

It can still be solved in an effective way with solvers like ...

TODO(gp): explain better, cite

**5.7. Reformulating the DaoCross problem.** In the DaoCross set-up, the price of the tokens involved in the swap is determined by an external oracle. This allows simplifying the general problem by applying the equivalence principle between orders and removing the non-linearity from the set of inequalities above.

In the case of each order  $o_i$ , the boolean condition  $c_i$  (which compares the actual price  $p^*$  and the limit order price  $p(o_i)$ ) is either verified or not (the solution does not need to consider both branches)

$$\begin{cases} 0 \leq q_{\text{base}}(o_i) \leq q_{\text{base}} & \text{if } c_i \\ q_{\text{base}}(o_i) = 0 & \text{if } \neg c_i \end{cases} \quad (3)$$

The problem then becomes a set of linear inequalities (in fact, if  $\neg c_i$  prevails, then order  $o_i$  is effectively removed from problem) that can be solved with various efficient methods (e.g., simplex-method).

TODO(gp): check

## 6. ALTERNATIVE SOLUTIONS FOR TOKEN SWAPPING

TODO(gp): Add colors: in green what's good, yellow, red

TODO(gp): Reorg, maybe from best to worst? DaoSwap, LOB, DaoCross, AMM

	<i>LOB</i>	<i>AMM</i>	<i>DaoSwap</i>	<i>DaoCross</i>
Custodial risk	Yes	No	No	No
Transparency	Low	High	High	Medium
Censorship resistance	No	Yes	Yes	No
Intrinsic efficiency	Medium	Low	High	Medium
Support for limit orders	Yes	No	Yes	Yes
Risk of predatory behaviors	High	High	Low	Medium
Impermanent loss	No	Yes	No	No
Multi-coin exchange support	Low	Medium	High	High

### 6.1. DaoSwap.

- Custodial risk: Low
  - Private keys and funds are always under the control of the users
- Transparency: High
  - Work is done on-chain and the off-chain computations are independently verifiable
- Censorship resistance: Yes
  - The application runs on-chain and even if the web front-end is attacked or disabled. The off-chain computation can be made robust using similar approaches to distributed oracles, or even ported on-chain if gas prices are not an issue (e.g., using Layer2 solutions)
- Intrinsic efficiency: High
  - DaoSwap relies on periodic auctions and it has been argued in previous literature that for small enough intervals (e.g., seconds) periodic auctions result in the same quality of continuous matching without allowing predatory behaviors
  - TODO(gp): Add a reference to Cramton paper
- Support for limit order: Yes
  - Limit orders are supported natively
- Risk of predatory behaviors: Low
  - TODO(gp):
- Impermanent loss: No
  - TODO(gp):
- Support for multi-coin swap: Yes
  - TODO(gp): DaoSwap/Cross pools all the liquidity in a single optimization problem

### 6.2. DaoCross.

- Custodial risk: Low
  - Same as DaoSwap
- Transparency: Medium
  - DaoCross relies on price discovery carried out on a lit exchange
- Censorship resistance: No
  - The robustness of DaoCross is the same as the lit exchange that it relies on. See comments on the corresponding topic for LOB
- Intrinsic efficiency: Medium
  - The efficiency in matching trades is the same as the limit-order book. See comments on the corresponding topic for LOB
- Support for limit order: Yes
  - Same as DaoSwap
- Risk of predatory behaviors:
  - TODO(gp)

- Impermanent loss: No
  - TODO(gp)
- Support for multi-coin swap:
  - TODO(gp): DaoSwap/Cross pools all the liquidity in a single optimization problem

### 6.3. Limit Order Book (LOB).

- Custodial risk: High
  - Exchanges own user private keys and funds, creating issues with fraud (e.g., FTX) and hacking (Mt. Gox)
- Transparency: Low
  - Many centralized exchanges can easily exaggerate volume (e.g., through wash trading) to attract ICOs and liquidity
  - There is no guarantee on the
- Censorship resistance: No
  - Users need to link their bank account to CEX and transfer funds, which can take 2-3 days and be subjected to Know-Your-Customer (KYC) and AML (Anti money laundering policies)
  - Governments can censor, interfere with, or even sever the connection between traditional and decentralized finance (see recent interventions of SEC against Coinbase and Binance, and FDIC rescue of Silvergate Bank, Signature Bank)
- Intrinsic efficiency: Medium
  - Limit order books are considered efficient in matching continuously trades, although recently researchers have criticized their time-continuous-time nature as source of latency arbitrage (TODO(gp): Add ref to Compton, IEX, Flash Boys)
  - The same token can be traded on different exchanges causing liquidity and price discovery to be fragmented with detriment on market quality
- Support for limit order: Yes
  - Limit order books naturally support limit orders from users
  - Risk of predatory behaviors: High
  - LOB operators often welcome and incentivize high-frequency traders as way to increase trading revenues and liquidity, at expense of predatory tactics (e.g., front-running, latency arbitrage, spoofing, sniping) which arm
- Impermanent loss: No
  - LOB
- Support for multi-coin swap: No

### 6.4. Automatic Market Makers (AMM).

- Custodial risk: Low
  - Private keys and the
- Transparency:
- Censorship risk:
- Intrinsic efficiency: Low
  - Arbitrageurs are needed to keep liquidity pool and prices in alignment with other centers for price discovery
- Support for limit order: No
  - Uniswap V3 doesn't support limit orders directly but only in terms of
- Risk of predatory behaviors: High
  - TODO(gp):
- Impermanent loss: Yes

- Uniswap is affected by impermanent loss because liquidity providers and liquidity takers are in general different users operating at different time scales
- Support for multi-coin swap: Medium
  - Uniswap requires multiple swaps and fees for arbitrary tokens

## REFERENCES

- [1] Hayden Adams, *Uniswap Whitepaper* (2018), available at <https://hackmd.io/s/HJ9jLsfTz>.
- [2] Hayden Adams, Noah Zinsmeister, and Dan Robinson, *Uniswap v2 Core* (March 2020), available at <https://uniswap.org/whitepaper.pdf>.
- [3] Hayden Adams, Noah Zinsmeister, Moody Salem, River Keefer, and Dan Robinson, *Uniswap v3 Core* (March 2021), available at <https://uniswap.org/whitepaper-v3.pdf>.
- [4] Alejandro Bernales, Daniel Ladley, Evangelos Litos, and Marcela Valenzuela, *Alternative Execution Priority Rules in Dark Pools* (July 22, 2022), available at [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=4169352](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4169352).
- [5] Jean-Philippe Bouchaud, Julius Bonart, Jonathan Donier, and Martin Gould, *Trades, Quotes and Prices: Financial Markets Under the Microscope*, Cambridge University Press, 2018.
- [6] CME Group, *2023 FX Product Guide*, available at <https://www.cmegroup.com/trading/fx/files/fx-product-guide-2023-us.pdf>.
- [7] ———, *Euro FX Futures - Contract Specs*, available at <https://www.cmegroup.com/markets/fx/g10/euro-fx.contractSpecs.html>.
- [8] Robin Hanson, *Combinatorial Information Market Design*, *Information Systems Frontiers* **5** (2003), 107-119, available at <https://doi.org/10.1023/A:1022058209073>.
- [9] Jason Milionis, Ciamac C. Moallemi, Tim Roughgarden, and Anthony Lee Zhang, *Automated Market Making and Loss-Versus-Rebalancing* (2022), available at <https://doi.org/10.48550/arXiv.2208.06046>.
- [10] Morgan Stanley, *Morgan Stanley Dark Pools*, available at <https://www.morganstanley.com/disclosures/morgan-stanley-dark-pools>.
- [11] ———, *MS Pool ATS-N Filings*, available at <https://www.sec.gov/cgi-bin/browse-edgar?action=getcompany&filenum=013-00117>.
- [12] *Walrasian auction*, available at [https://en.wikipedia.org/wiki/Walrasian\\_auction](https://en.wikipedia.org/wiki/Walrasian_auction).
- [13] Mu Xia, Jan Stallaert, and Andrew B. Whinston, *Solving the combinatorial double auction problem*, *Journal of Operation Research* **164** (2005), 239-251, available at <https://www.sciencedirect.com/science/article/abs/pii/S0377221703008981>.